

# WENXUAN ZHOU

Email: [zhouwenx@usc.edu](mailto:zhouwenx@usc.edu) Homepage: <https://wzhouad.github.io/>

## EDUCATION

**University of Southern California, Department of Computer Science**  
Ph.D. candidate in Computer Science  
Advisor: Muhao Chen

Los Angeles, CA  
Sept. 2018 – Present

**Hong Kong University of Science and Technology, Department of Computer Science**  
B.Sc. in Computer Science and B.Sc. in Applied Mathematics

Kowloon, Hong Kong  
Sept. 2014 – May 2018

**University of Illinois at Urbana-Champaign**  
Exchange Student in Computer Science

Champaign, IL  
Jan. 2017 – May 2017

## HONORS AND AWARDS

<b>Best Paper Honorable Mention Award, WWW 2020</b>	2020
<b>USC Annenberg Fellowship</b>	2018 - 2019
<b>Continuing Undergraduate Scholarship for Academic Excellence</b>	2015 - 2017
<b>Dean's List, School of Engineering, HKUST</b>	2014 - 2017
<b>School of Engineering Scholarship, HKUST</b>	2014

## PUBLICATIONS

1. **Wenxuan Zhou**, Sheng Zhang, Tristan Naumann, Muhao Chen, Hoifung Poon. Continual Contrastive Finetuning Improves Low-Resource Relation Extraction. In the 61st Annual Meeting of the Association for Computational Linguistics (**ACL**), 2023.
2. Keming Lu, I-Hung Hsu, **Wenxuan Zhou**, Mingyu Derek Ma, Muhao Chen. Multi-hop Evidence Retrieval for Cross-document Relation Extraction. In the 61st Annual Meeting of the Association for Computational Linguistics (**ACL**) - Findings, 2023.
3. Fei Wang\*, James Y. Huang\*, Tianyi Yan, **Wenxuan Zhou**, Muhao Chen. Robust Natural Language Understanding with Residual Attention Debiasing. In the 61st Annual Meeting of the Association for Computational Linguistics (**ACL**) - Findings, 2023.
4. Xiacong Yang, James Y. Huang, **Wenxuan Zhou**, Muhao Chen. Parameter-Efficient Tuning with Special Token Adaptation. In 17th Conference of the European Chapter of the Association for Computational Linguistics (**EACL**), 2023.
5. **Wenxuan Zhou**, Fangyu Liu, Huan Zhang, Muhao Chen. Sharpness-Aware Minimization with Dynamic Reweighting. In the 37th Conference on Empirical Methods in Natural Language Processing (**EMNLP**) - Findings, 2022.
6. Keming Lu, I-Hung Hsu, **Wenxuan Zhou**, Mingyu Derek Ma, Muhao Chen. Summarization as Indirect Supervision for Relation Extraction. In the 37th Conference on Empirical Methods in Natural Language Processing (**EMNLP**) - Findings, 2022.
7. **Wenxuan Zhou**, Muhao Chen. An Improved Baseline for Sentence-Level Relation Extraction. In the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (**AAACL**), 2022. (short paper, **Oral**)
8. **Wenxuan Zhou**, Qiang Ning, Heba Elfardy, Kevin Small, Muhao Chen. Answer Consolidation: Formulation and Benchmarking. In the 20th Annual Conference of the North American Chapter of the Association for Computational Linguistics (**NAACL**), 2022. (**Oral**)
9. Yiwei Wang, Muhao Chen, **Wenxuan Zhou**, Yujun Cai, Yuxuan Liang, Dayiheng Liu, Baosong Yang, Juncheng Liu, Bryan Hooi. Should We Rely on Entity Mentions for Relation Extraction? Debiasing Relation Extraction with Counterfactual Analysis. In the 20th Annual Conference of the North American Chapter of the Association for Computational Linguistics (**NAACL**), 2022. (**Oral**)
10. Yiwei Wang, Muhao Chen, **Wenxuan Zhou**, Yujun Cai, Yuxuan Liang, Bryan Hooi. GraphCache: Message Passing as Caching for Sentence-Level Relation Extraction. In the 20th Annual Conference of the North American Chapter of the Association for Computational Linguistics (**NAACL**) - Findings, 2022.
11. **Wenxuan Zhou**\*, Fangyu Liu\*, Ivan Vulić, Nigel Collier, Muhao Chen. Prix-LM: Pretraining for Multilingual Knowledge Base Construction. In the 60th Annual Meeting of the Association for Computational Linguistics (**ACL**), 2022. (**Oral**)
12. **Wenxuan Zhou**, Fangyu Liu, Muhao Chen. Contrastive Out-of-Distribution Detection for Pretrained Transformers. In the 36th Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2021. (**Oral**)

13. **Wenxuan Zhou**, Muhao Chen. Learning from Noisy Labels for Entity-Centric Information Extraction. In the 36th Conference on Empirical Methods in Natural Language Processing (**EMNLP**), 2021.
14. **Wenxuan Zhou**, Kevin Huang, Tengyu Ma, Jing Huang. Document-Level Relation Extraction with Adaptive Thresholding and Localized Context Pooling. In the 35th AAAI Conference on Artificial Intelligence (**AAAI**), 2021.
15. **Wenxuan Zhou**, Bill Yuchen Lin, Xiang Ren. IsoBN: Fine-Tuning BERT with Isotropic Batch Normalization. In the 35th AAAI Conference on Artificial Intelligence (**AAAI**), 2021.
16. **Wenxuan Zhou**, Hongtao Lin, Bill Yuchen Lin, Ziqi Wang, Junyi Du, Leonardo Neves, Xiang Ren. NERO: A Neural Rule Grounding Framework for Label-Efficient Relation Extraction. In the ACM Web Conference (**WWW**), 2020. (**Best Paper Honorable Mention Award**)
17. Ziqi Wang\*, Yujia Qin\*, **Wenxuan Zhou**, Jun Yan, Qinyuan Ye, Leonardo Neves, Zhiyuan Liu, Xiang Ren. Learning to Annotate: Modularizing Data Augmentation for Text Classifiers with Natural Language Explanations. In the 8th International Conference on Learning Representations (**ICLR**), 2020.
18. Ziqian Zeng, **Wenxuan Zhou**, Xin Liu, Yangqiu Song. A Variational Approach to Weakly Supervised Document-Level Multi-Aspect Sentiment Classification. In the 17th Annual Conference of the North American Chapter of the Association for Computational Linguistics (**NAACL**), 2019.

## PREPRINTS

---

1. **Wenxuan Zhou**, Sheng Zhang, Hoifung Poon, Muhao Chen. Context-faithful Prompting for Large Language Models. Arxiv.
2. Tianqing Fang, **Wenxuan Zhou**, Fangyu Liu, Hongming Zhang, Yangqiu Song, Muhao Chen. On-the-fly Denoising for Data Augmentation in Natural Language Understanding. Arxiv.

## WORK EXPERIENCE

---

- Microsoft Research, Redmond, WA** May 2022 – Aug. 2022
- **Research Intern**, *Continual Contrastive Finetuning Improves Low-Resource Relation Extraction* (in submission).
  - Mentors: Sheng Zhang, Tristan Naumann, and Hoifung Poon
- Amazon Alexa Search, Seattle, WA** May 2021 – Aug. 2021
- **Applied Scientist Intern**, *Answer Consolidation: Formulation and Benchmarking* [5].
  - Mentors: Qiang Ning, Heba Elfardy, and Kevin Small
- JD AI Research, Mountain View, CA** May 2020 – Aug. 2020
- **Research Intern**, *Document-Level Relation Extraction with Adaptive Thresholding and Localized Context Pooling* [11]
  - Mentors: Tengyu Ma and Jing Huang
- Bytedance AI Lab, Beijing, China** May 2019 – July 2019
- **Research Intern**, a universal method for improving fine-tuning performance on PTLMs.
  - Mentor: Hang Li
- Hong Kong University of Science and Technology, Kowloon, Hong Kong** July 2017 – Feb. 2018
- **Research Intern**, *A Variational Approach to Weakly Supervised Document-Level Multi-Aspect Sentiment Classification* [15]
  - Mentor: Yangqiu Song

## PROJECTS

---

- Reimplementation of R-Net** 500+ ★
- A Tensorflow reimplementation of the paper “R-Net: Machine Reading Comprehension with Self-Matching Networks”. The first Github repository that succeeds in reproducing the reported results.

## SERVICE

---

**PC Member:** LLD 2019, DeepLo 2019 / 2022, Repl4NLP 2019, EMNLP 2019 – 2022, AAAI 2020 – 2023, ACL 2020 - 2023, AACL 2020 / 2022, NLPCC 2021 / 2022, ARR 2021 - 2023, TNNLS, NCAA